

Solidity Compiler Tools Design Document

Killian Rutherford (KRR2125)

Raphael Norwitz (RSN2117)

Jiayang Li (JL4305)

- 1. Problem**
- 2. FTime-Report Tool Design**
- 3. Optimisation within Compiler Source Code**
- 4. Parallelisation**
- 5. Parallelisation Tool**
- 6. Future Work/Version 1.2**
- 7. References**

1. Problem

Compiling large solidity files, particularly large solidity libraries with multiple file, can be quite time consuming with the current compiler (some cases taking several 10's of seconds). The tool we developed would enable users and developers of the solidity compiler [1] to profile pieces of code/functions within the compiler to see which areas were consuming a lot of time and acting as bottlenecks.

2. FTime-Report Tool Design

(See Tutorial for Introduction and User Manual for in detail about every element)

All our code is located on our github page <https://github.com/raphael-s-norwitz/solidity> [2].

Our initial design option was to produce a lightweight solution that would allow flexibility on the user/developer side in terms of what code to profile. This initial design is produced on the dev_tool branch. These flags would not alter any core compiler functionality, but provide a function call trace and time readings for each function call in the call stack.

Hence we opted for a stack-based approach for simplicity and efficiency, with TimeNode classes to hold information of begin and end times when the user pushes/pops onto the stack, names, and vector of TimeNode representing its children. We referred to gcc's implementation of --ftime-report [3] when designing our tool.

Chrono::high_resolution_clock was utilised for time measurements given it's accuracy and what we learnt in class.

Initially we wanted to have the time profile code be implemented as a base class which other classes and functions might be able to inherit from to print out information, but this was too challenging/not possible.

We decided that if a developer wants to benchmark a function or segment of code he/she will have to `#include <libsolidity/interface/ftime.h>` in the file; add `t_stack.push(string)` at the beginning of the function/code, and `t_stack.pop()` at the end. This adds flexibility on the user's behalf, as he/she can decide what is to be tested, how often within a function/class to test, and give appropriate string descriptions as to what is being tested.

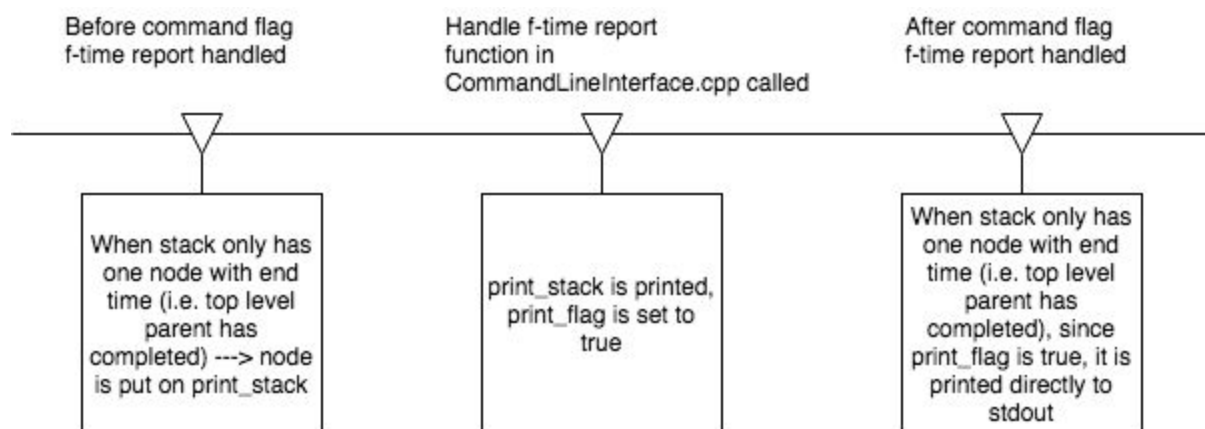
TimeNodeStack holds these TimeNodes; so that whenever a user pushes to this stack given a string argument, a new TimeNode is created with name equal to argument, begin time at the time the user pushes, and put on a stack represented as a vector.

When a user pops, if the stack has multiple elements, the top of the stack is popped (end time of this node is set) and this TimeNode is added to the stack's new top item's children list. This makes sense from a design perspective as no pointers are necessary for a TimeNode to point to its parent. If the tool is used correctly (no errors are thrown), and if the stack has more than one element, it means that the user has nested push statements (whether in the same function call or in functions called by higher-level functions).

If a user pops when the stack has only one element, it means the end of the recursion in terms of pops and pushes for the highest level call has occurred, and this can be either printed if it happened after the compiler handled the `--ftime-report` flag, or added to the `print_stack` so that the compiler can decide whether to print the report when it checks the value of the `--ftime-report` flag.

Having the user write `t_stack.push(string)` at the beginning and `pop()` at the end allows flexibility on the user's part in terms of being able to profile multiple times and multiple positions within the same function block [2].

`libsolidity/interface/FTime.h` defines a global `TimeNodeStack`, which the user can push to/pop from. It is necessary to have a global variable if the user wishes to profile code within the compiler that gets executed before the flag `--ftime-report` is handled. This comes hand in hand with the `print_stack` mentioned before. Since the developer can choose to benchmark code executed before or after the `--ftime-report` flag check, in order for the flag to work properly (i.e. data can only be printed out if the flag is specified), another `print_stack` needs to be created to store everything that is benchmarked before the run-time argument is treated. Then, another flag is turned on to print everything that is benchmarked after the run-time command line argument parsing code has been executed. (See diagram below for clarification).



However, users can also define multiple `t_stacks` if they want (this came in handy in testing), and these can become very useful when linked with `TimeNodeWrappers` (discussed below).

We also added a `printString` function to print the `print_stack` to a stringstream rather than dumping to `stdout`; to enable test cases to be handled much more easily.

We decided to only measure in microseconds because we expect that any performance hit in the solidity compiler can be measured on a reasonable scale in microseconds. Also, we hoped to build a handy tool for users and therefore wish to eliminate any unnecessary the parameters from our API.

We also incorporate the tree optional flag “time-report-no-tree” to give users more flexibility. If they do not want to print out in tree form, they can specify the flag and have the plain form.

We noticed that in many functions where there may be several early returns, we needed to put `pop()` before each of them. This can be time-consuming and error-prone since a user may not notice where an early return happens. Hence we decided to incorporate RAIL into the tool to handle resources on the stack more intelligently (found on `dev_tool_wrapper` branch).

We created a `TimeNodeWrapper` class which would minimize the number of `pop()` calls needed to profile a complicated block. Using RAIL principles learnt in class, `TimeNodeWrapper` object is instantiated with a `given_stack` and a given time node name. In the constructor of `TimeNodeWrapper` we call `push` on the `given_stack`, and set `TimeNodeWrapper`’s private stack the same as the `given_stack` so that we can refer to `given_stack` when destructing the object. The `TimeNodeWrapper` also has a boolean member called `popped` that is `false` by default. If the user manually pops by calling `TimeNodeWrapper`’s `pop()` API, the `popped` flag will be set `true`. When the object is being destructed, the destructor will check the value of `popped` and it will only `pop` from the `given_stack` if the boolean value is `false`. We believe this design offers users greater flexibility when profiling function blocks in the solidity compiler code [2].

Further, the destructor pops from the stack, it will check whether the name returned from the `pop` operation matches the object’s name. If the names don’t match, it will throw an error. We are working under the assumption that objects are destructed in the reverse order of construction; as one would expect from how memory stack allocation works. This allows the user to define multiple `TimeNodeWrapper` objects within a block of code or function in order to profile different sections (such as loops for instance). We decided that the constructor to a `TimeNodeWrapper` must take the stack and a string. We thought about writing another constructor which does not require a string argument to name the Node, and would instead take the function name within which it was called. However, this would make it harder to define several `TimeNodeWrappers` within the same function block, (an extra count variable to keep track may have had to be added); and if two functions in different namespaces had the same name, it would be hard to keep track of which was which. We hence thought it best in terms of simplicity and flexibility to keep the string argument in the constructor [2].

3. Optimisation within Compiler Source Code

In the beginning, we noticed that running the test script was taking a lot longer on some machines, we investigated, running on Mac vs. linux, but turned out to be optimisation time differences whether building with make or ./scripts/build.sh. The latter built the 'Release' version solc while the former built the 'Debug' version [1].

After creation of version 1 of the tool, we began to use it to see how we could identify the slow passes of the compiling process and possibly optimise on various things.

Looking at section 3.a of the User Manual, we can see the results from further testing the optimise() functions within the compiler. We see that the applyMethods loop is taking the most time. (code seen on dev_optimise_profile branch).

The bottleneck seems to be that each of these methods (PushPop, OpPop, etc..) is applied sequentially. We attempted to write out the applyMethods iteratively, calling each Method instead of recursively calling a variadic template function, this did not seem to improve the timing.

Further attempts to take out the template metaprogramming were unsuccessful. Their approach of using static functions avoids may constructor calls, while complicated, seems to be the faster than any naive approach we can come up with.

We also considered calling these variety of methods asynchronous, but we realized it would be complicated and involve several locks most likely since the optimiserState would be altered, and actually these methods may need to be called in a specific order.

For results in investigations into the analyse pass of the compiler, see the User Manual section 3b. Attempts to shave off time to refine the source code were unsuccessful.

4. Parallelisation

At the moment, solidity lacks any kind of parallelism, and this seemed like the most promising way to get speedups.

As discussed before, once the ./scripts/build.sh script is run, and the solidity executable is produced (solc); we compile a solidity source file (.sol file) by running solc Coin.sol for example, along with a variety of flags (bin, optimize, ignore-missing), including path to produce the binary file (-o dir). The solc executable can take in a list of multiple source code files (and all need to be together if dependencies amongst files exist).

There are two obvious courses to reduce compile-time through parallelisation when dealing with multiple file libraries.

1. Higher-level parallelisation: splitting files across different processors, running compiler executables on each processor.
2. Lower-level parallelisation: parallelization within the compiler source code

One would hypothesise that the first method above would only be beneficial for compiling large numbers of solidity files. From above tool measurements, there is a fixed time cost for the compiler to be instantiated, and thus by doing so on several processors, a certain number of files will be necessary to offset this instantiation cost and make splitting compiling tasks among several processors worthwhile.

Further, it became clear from the onset that option number 1 would only be possible for solidity files which did not have mixed dependencies (if random splitting of the files was applied). In the best case, each processor would have to be self contained sets of dependencies, but this may be hard to identify without some sort of parsing through source code to see which files are included where (probably too hard to do non-manually).

If files are unrelated and have redefinition of contracts, it will pose issues and potential errors when compiling them together.

A lot of libraries which have inter-dependencies among files would hence make it difficult to apply the first optimisation technique.

Initially, a simple shell script was written to split the compilation across different processors.

The number of files given a certain number of processors on the machine could be calculated optimally using UNIX commands. Simple division and modulo remainder can split everything optimally [2].

E.g. given 10 solidity files, A.sol - J.sol, and 4 processors, we would compare running:

```
solc A.sol B.sol C.sol D.sol E.sol F.sol G.sol H.sol I.sol J.sol
```

(on 1 processor)

Vs.

```
CPU1: solc A.sol B.sol C.sol
```

```
CPU2: solc D.sol E.sol F.sol
```

```
CPU3: solc G.sol H.sol
```

```
CPU4: solc I.sol J.sol
```

Initial tests of with vs. without parallelisation showed some potential for improvement. Although not taking into account dependencies correctly in some cases (i.e. checking for correct

compilation), we saw before parallelisation on a google instance with n1-standard-4 (4 vCPUs, 15 GB memory), over 10 runs an average compilation time of 4.81 seconds with stdev 0.014, whilst with using parallelisation- 4.23 average with stdev 0.05 seconds, giving an approximate 12% speedup. However, we noticed that using a bash script was not the best idea for portability across systems and the potential overhead it introduced. Additionally, no checks were done to make sure everything was compiled correctly [2].

Hence the switch was made to use `async` in a `cpp` file instead to document and measure everything properly (`dev_optimise` branch).

Using `system()` to call the compiler with relevant flags, and thread library, all files were split in a vector according to CPU and run on different threads using `async` (with the `launch::async` flag set), and `get` used to obtain the future (stored within a vector of futures).

Although there are issues with using `system()` such as portability (<https://stackoverflow.com/questions/8086071/are-system-calls-evil>) [4], we needed to be able to call the `solc` executable, and we assumed since the `solc` executable is portable, we were within right to use `system()`.

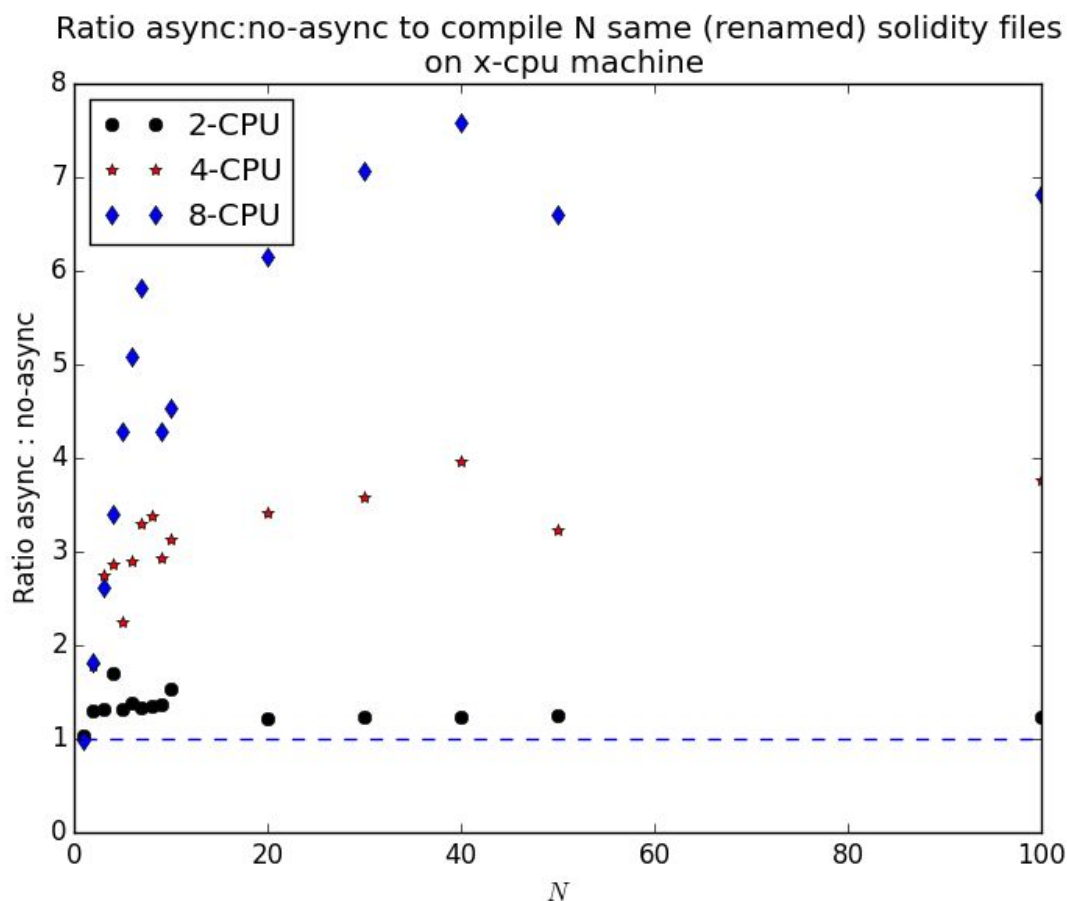
We ran into dependency issues with the test libraries provided by the solidity community, hence to determine whether parallelisation was helping, we needed to decide upon a dataset in which we know all files could be compiled independent of one another.

It may be possible to figure out which source files need to be run together given their dependencies, but for larger libraries (unless some automated tool existed), manual selection and reading header files would become very impractical. As mentioned in further work at the end, it may be possible to build a topological sort of the files in terms of when to compile them, by checking in the source files which ones import which.

The first thing we thought of was to copy a file numerous times within a library, however this would give compilation/dependency errors; hence we moved on to selecting a random contract, and copying it a certain number of times, while altering the name of the contract classes and functions to prevent the compiler from optimizing and predicting / cache hits.

This was done with `BlindAuction.sol` file, which contains both a `BlindAuction` contract and a purchase contract. (see <https://solidity.readthedocs.io/en/v0.4.23/solidity-by-example.html> for the files) [5].

Time ratios were averaged over 3 runs and results over different number of google instance n1-standard vCPU are shown. (n1-standard-2 (2 vCPUs, 7.5 GB memory) ; n1-standard-4 (4 vCPUs, 15 GB memory) ; n1-standard-8 (8 vCPUs, 30 GB memory))



We can see pretty quickly the benefits of parallelisation and having multiple CPUs available. As soon as N goes beyond 2, we see ratio increases, most likely due to the fact that this particular solidity file is large enough so that compiling all of it offsets the fixed cost of starting another executable on another processor. For larger number of files we see a larger separation and approximate doubling of ratio with doubling of number of CPUs.

We also wanted to test out parallelisation on a more realistic data set (i.e. not just copies of a specific file) to see what could be a potential “real-world” speedup to benchmark and to avoid any thread taking advantage of cache memory. We were pointed by the solidity development team towards <https://github.com/allenday/github-solidity-all> [6], which contains all solidity files on github in one place.

Rather than selecting x number of files and manually checking if they could be compiled/split across different processors, we automated this with C++ code and the `system()` function call.

Initially, we decided to run through different files in the github solidity all repo + the test files provided by the solidity community. We would compile the first file, using try-catch to catch any

exceptions from the compilation. If any errors were caught the file would be discarded and the next one tried. We would keep on adding files this way until we obtained the number of files wanted.

We quickly realized that `system()` was not throwing exceptions, but would return an integer code instead; so our first check to see if the files were being compiled together correctly (before attempting to parallelise) would be to check if the return code was successful (0), then add the file to the list, otherwise discard and move on to the next file.

However, even only checking for the return value was not enough, as in some cases a correct return value was returned, but the compiler did not actually compile (specific error was the respective binary file was not generated as a compiler version was not specified in a particular solidity file). Hence, the check for correctness would have to be whether a correct binary file was produced (which was achieved through the `--bin` flag). With these, we were able to obtain 100 files from the general solidity files on github [6], which were able to compile independently and move forward with parallelising.

100 files:

8894 lines; 33,362 words; 294,515 characters

```
73 187 1955 ./test_files/test_a7dbaa67d524dff15145be4b36f6d41f7b245e8221fc4004e05677268dd33a52.sol
24 78 454 ./test_files/test_f1ddb3ec5142f65be5e1e26a6a476aa8f06d2c83e9647968d2f27041a9b4c980.sol
73 187 1955 ./test_files/test_ce630e9d43f8716989b99d88bf00b32fb101a429574ad3ed92df0cfb64f979a2.sol
53 197 1106 ./test_files/test_cf7f78c987171edefc69321edf6e5b9f8ec2e4b0dbfd0c0dd43367a034c9f16c.sol
13 73 700 ./test_files/test_8da06f832f837a554e151371368750f2b39a2f749073e7261d7ad02dc61a10ee.sol
78 302 2521 ./test_files/test_beb72a1741f79c88d654d63439099a12e8ae8ea39330f8a3a31ca4bfac78ca59.sol
14 53 445 ./test_files/test_a9cf1d9073a8a58ca6044a1720a93a69020ea80fab3f5169192630590707d593.sol
361 1336 12385 ./test_files/test_41145dd32e305cda546cada00805fe647a7db9dcadee99ed4a107f075144b59b.sol
9 23 201 ./test_files/test_84025419be75907457e68902da17c2ba380ebb86aeda73a1701501d03bcb319e.sol
98 370 2341 ./test_files/test_f05adf65bd6a71de586500a74b8e07a9f4a70625e90c55e93ebb5e400ddfdb41.sol
6 27 158 ./test_files/test_2d4cd83f17573d726a08d119313f6aaae9ba47f8439e6247e06057e295452ba.sol
55 178 1859 ./test_files/test_ced95337821dd3febb160aedcc50ad760eb75971924657e14858606aaf1fdf29.sol
31 53 506 ./test_files/test_eff0a9670ebf709ee0b054bd3dd7fde373314e4fea9348803eccab5cccd3dd80.sol
102 450 3487 ./test_files/test_f05dd9bacc2604127cf72d2038b5c42fac0a832d44b210f129f138cb400f2b5.sol
13 35 281 ./test_files/test_e2fd2d93f8741884e1b9d0808d0724a13c9082da5e247ed6f8f46e26dad60719.sol
281 856 8502 ./test_files/test_180ac4cf9e021cdd42c556aa203e951aab4a8a0fb52a7b436fa9b5977f7e998a.sol
109 500 3923 ./test_files/test_729c48d31247252239590ba80a77aa63be3ef9feeddc1aef92ba095f07371033.sol
370 1102 11515 ./test_files/test_c4c8cf07bc07d54727a24fa9ab84e7353637274d213406ba663bf1848ad696e6.sol
19 45 500 ./test_files/test_27548d90ec169921a0c7ca596b266088059e2bb91976225ecab57be131e002a7.sol
23 82 701 ./test_files/test_caf1d08f8255ca913796af6a5dcf3198b250ea8b864c7293cb815ade0841b0f7.sol
7 20 101 ./test_files/test_1a92c9e1ad65fb0e2a67ed5f7e71cb5f2f640e701f5ca69c56abc3b175c6a111.sol
79 195 2367 ./test_files/test_ebfcda9f83b927375ead8de9a0bc57dc23a9326cd204ea24797b62d80b887b99.sol
9 23 202 ./test_files/test_aa49ec7f336c6fea1204a2a4152c63dbd56ddeca0c1a348bbd51783dd671ed0e.sol
15 79 747 ./test_files/test_c5eeede822366b11933471f25adbbe999e1c3e993ea881af9ca7e63c79b90df9.sol
26 84 596 ./test_files/test_cc48ccd774362b86a68ac59b272eaf32501e4aa87c93d5a38c6df55964cc854a.sol
800 4066 32437 ./test_files/test_b32af89dda8f16ada02ed5a3dd800f3fe17230bc845c1483d0f4a2ed7faa78aa.sol
708 2238 24511 ./test_files/test_b1d17622bfd063291ef15adf99fa8860ac01f994ac631fada2660be73152d1ae.sol
15 54 392 ./test_files/test_a9158c87554cd9f67dee4e450f01b3335cd86eb91ee546b5855d48a574d12800.sol
60 229 1719 ./test_files/test_6258bfb64990fa7fe594dda3a0eafcf12cd3d9113eaf5a9e5d07589cc0d47b2e.sol
```

592 2984 24208 ./test_files/test_881cbfc5553c77da2fbbee2086dc100069d058980405477e0327da491b0ff320.sol
32 99 877 ./test_files/test_f07594dfb6630bf2bc3412854dc6dcb7be1aa82889431e073086ceefefc1b3f.sol
42 92 704 ./test_files/test_55d6d7a52cac81df9056bc530b703b7aa6a9d6f319ab1aea3bb3920cb9975cba.sol
27 56 504 ./test_files/test_94cd7a46e80fbfb80989cb9b2d712e7f3606683f23762e534ee4b558085597e6.sol
20 35 339 ./test_files/test_a389cb0ed307d0fd0f0496684d8b0acf01a8be2603f32fb4f528a0ef146328cd1.sol
18 51 398 ./test_files/test_d23f90ac6b6d7c5654023b79bbd39281fbb741dfee8bb6ec3385ab30f542530.sol
6 19 146 ./test_files/test_0e08e1aa8ac483d1d2683286aac69dc0229ad0ea5955fcee03ba27e27e14e94.sol
17 32 270 ./test_files/test_2e67c6ba377ef76a380cb76d53ff5b254d123c649825f27b2c7dc2b7bdd73a30.sol
23 48 450 ./test_files/test_b0eac34cd208612102c3aa9d5228bc9cc5b8101313bbeb59f692955f0b60fd719.sol
90 335 2959 ./test_files/test_c240ac4f115b4790eedc59a555d6c75640a9ec635ee94b9a252292bed07117ca.sol
172 429 4370 ./test_files/test_27b57c491d76d89275c92bfe517c7c466f9c5baaa420d585b0780ebe88aa5032.sol
11 20 174 ./test_files/test_53f3679909545e5cb34d838087dd16ffd75f8907cdf352168b2b2dbb08862f1d.sol
260 946 10439 ./test_files/test_00fef28dbf32dabeeedda75bed666550761cb4a58fb49bdb2779bc5c96fe6fb7.sol
53 213 1707 ./test_files/test_24b5acc4187857057a3b3acc67e8d375c1e99f85227d028ebe3143a24e02ca5e.sol
20 35 381 ./test_files/test_2600be00203e59b5ea26b972e18058bd3e3e3b51a20b9caf77b295466a96cfcd.sol
91 343 3394 ./test_files/test_bae4d162fd7d404dbbb62da7c6f00cacac4e95ccb283328f485146141728376e.sol
23 52 515 ./test_files/test_1073beba092f03f5d3425f42d7b5b1df41f37f9949cb339126a12105d97df320.sol
2 4 15 ./test_files/test_69dbaef3b2304c096a355e9871d6bf70e654b22b97c617210bb6ccd270b874c5.sol
220 702 5220 ./test_files/test_de24c86b1a6b9fc3158132ba43f6710091fa1d6637dc335514d84f7bd0973307.sol
56 133 1800 ./test_files/test_cc279795bcb1be51ca7379cbcd096b82301f3cd0c6c80b05a6644b90d86130776.sol
23 49 493 ./test_files/test_d1ba9cc38ae66d7c7b6f2248262205cbe5d76452608f92e0428c2174c7a59376.sol
99 516 3949 ./test_files/test_c5c14d4ce17800de3fa979a48306518a82c35acb8240e426c586d447c2db1d65.sol
8 32 463 ./test_files/test_dcd0ee37c687155107d55c5ea7c64a1b30ec8329beb139fd1315f215938be06.sol
9 20 140 ./test_files/test_30b8209f11035d7ab60d8335e970f58fc7a592c2780497e9a3d07e3a08b3c277.sol
17 56 631 ./test_files/test_ee27b9d841c89e58b670db6d7c2096b0cebd621e4224293caa2ac3074ea678de.sol
31 125 866 ./test_files/test_b2a1e02d34d25a72ad5a005795d9305be4dd6009b6c31122fbc8e0963536489.sol
24 49 495 ./test_files/test_079bcf654c9fb9acd811faf181ffa8dc2921a059dd430135fc474e8f427422d9.sol
20 56 582 ./test_files/test_b727e7cdc55f3caef68c6ba823f9047752573090ece5d5a412946c12ed9e53e.sol
21 67 483 ./test_files/test_d09447013d256b523b67b951ccb5a5719b20e38f9e47a6a10bf28d1e5de31104.sol
83 183 1978 ./test_files/test_75e04c1d769f47687a3b6db1f040944cff9408be0d5cc8e42038f2aca05f5af0.sol
26 45 407 ./test_files/test_a1bcc24b41baa69febd36c222d0104c562c6dda12e1175d10168cf3e3aedb73d.sol
364 1098 11381 ./test_files/test_1dd8c290ab7230fc1267d8ee3856f6578d3e2235502aacad51f7a21f1594408e.sol
152 678 5653 ./test_files/test_391ed7631d9a1183c6d6293defce1925053bc0f056e7650a762122003e31a16d.sol
269 773 9115 ./test_files/test_7166048567b924db185cd3804f51f6f3969af314255f8494c8cf8338e06288cf.sol
4 12 107 ./test_files/test_5f913b422c3a08a1d37d6264c43d5faaa361cd2ff14669ee0b7231d618660c7a.sol
103 274 2196 ./test_files/test_ff599bc711526922add29d87d049be740bf4638a3cbadda297e462d721f85a92.sol
91 311 2476 ./test_files/test_543ac5f8897d242b076187b3d9a169b1dac4538f4229240f8819bc30b2952f49.sol
41 205 1441 ./test_files/test_af75ba8c2005cdabd2e31bf9cf0e6e510fc64c1f9bd2a621347d537f14caeffa0.sol
46 118 645 ./test_files/test_40533309e2d270c161b9d328f331d028883c319de13cce6dff878dcb9b9480e2.sol
13 33 321 ./test_files/test_14c7a213c9ee60b1e241b74e196892fd0675f46a8f4bcd3b27acbeea0bf7eca6.sol
713 2968 24859 ./test_files/test_d83f6fd57215fe8c796d4a8f8c08375b206fee9e9da061fb4d2861bc7feb3fa3.sol
45 100 861 ./test_files/test_145d369f014d451f55bc43acc7eb38d49423046447e14212f61984ba5daa7683.sol
366 1853 14384 ./test_files/test_f11a1bc2f65010799a1a82fc9ca3324b51667841eebf3e95129f54b7ede9e315.sol
17 53 408 ./test_files/test_80e6a9f1bc2db5c8cec14e462b886ba214c2fe09c748d720a764a6a59575cc3e.sol
13 53 341 ./test_files/test_94d476d127829e60e66b0ef454ebe732efcca46da3f27f4d0b54c995cb6bd31a.sol
23 49 493 ./test_files/test_8313e2d3233bffe2cdaac2213160eef056d25a6ad69dfb4aa8c93a78ed6276fd.sol
0 3 14 ./test_files/test_bc432a88a5ecee79a418334fedd1a53e35ad068d12ac8ee7a59605365b8ace15.sol
15 29 255 ./test_files/test_d4d2ec889a6558c5961ce550e9fddfc5863818adf4f493895f11c36d52d2627c.sol
14 76 720 ./test_files/test_fdca83ed60188aa46b3160cdccd28b118b9e56cf92dda047aa993325323393f3.sol
108 427 3920 ./test_files/test_3991bb702569121efda01c9ad3df8b8089a6e5fa17e0f4a30a4976e2858debae.sol
49 218 1215 ./test_files/test_ac816db87fad29b87c7da4a9c22da62b39d1f1fb808c06b68476ab0e63dc72db.sol
17 47 369 ./test_files/test_d17c2e430477347046df44382dfc766e1e0efd508cc320a34f91272cfe6ddecc.sol
140 792 6440 ./test_files/test_475316a006533f57931d2ec6efa3e6be030242a63989d77dcfdc8efb4fa60c49.sol
55 217 1803 ./test_files/test_c76bbdfdbde7969a1103bcc14d042a6f3389948fa42d0c1eea3201e24dceb9a.sol
4 18 121 ./test_files/test_4fc6fc1064eb624d65550ee72eeb43edba8b41966f9b657e99dbe2067376eeb.sol
18 76 647 ./test_files/test_eec0e7b013ac7adf7ecb18ff453bf187ef4f6978ff7293c79a840f38db77a243.sol
13 22 162 ./test_files/test_b3e5b2195f09b605e67c8016dabd8167ea904074debe0b1151869d245689d17c.sol

```


144 541 4602 ./test_files/test_d5ba0ba7d5a54e4c4af70c653eff6732f591bc7a0dd0c31c8d277731412d0569.sol
23 39 299 ./test_files/test_87b1ce31b14175f8a7f66bb5879edd12ff4fbca420ac83a8ebdee48d11f82de1.sol
13 33 268 ./test_files/test_27329c49aa09c3a886e4ddd8002ca61485f5f3624d353fb0a6d4bc22b911cfa.sol
25 82 668 ./test_files/test_e6889d6b48abf30a71a3cbb79a0e496521e2c11e7c2d342358f5b481f19bfe2d.sol
57 191 2466 ./test_files/test_0c1ea4a45c5036b405bfcd915e0a4f859c280eb53c18247f53b1bdc8f5999e46.sol
21 64 621 ./test_files/test_f4eaa60c512f93cd976c80ecc6a65ea57b3bef16de3ac37af926f650f0b7d0b3.sol
73 308 2071 ./test_files/test_25e04cd2828edeee0dec5fe14f5da908705ab7a82c510bb25c880adca7cfa62.sol
55 337 2508 ./test_files/test_b3fff7e6a3bff46da8a8493ebac5edc8dd9665ec8f4fe4f5953caab3f92f25b0.sol
14 27 225 ./test_files/test_8bf105fee7e33f095ec43ebfda22d3a2a133f72fd4ebb84beab8dda0cf6a26bf.sol
8 14 95 ./test_files/test_ceddcb74312958562b83418eb10f2df74f9be70c76688eed28c00c3729f9baf9e.sol
9 39 260 ./test_files/test_a201fac7f454b4b98880d79141bf8f23a53430a6620c4a3c4b300154ba0be338.sol
55 130 1429 ./test_files/test_336feb9082de244fa180ce831404e8426e445dc61b3e15ce42ef054130fff292.sol
27 38 351 ./test_files/test_518fc11b394e60cf0f3ea7e5c397526dbf87e7b6c299af7f14ef3f838f0f94ee.sol
17 68 381 ./test_files/test_f805a68f0675fb683edb12a6bfa19d777a1b75c6d4bcc6e6502b41db3b511cc8.sol
8894 33362 294515 total

```

Finally, to test if the parallelising compilation produced the same binaries as the original solidity compiler, we ran a python script to check if the binaries compiled on 1 CPU vs. multi were exactly the same. In all of our tests, if two files generated by the above 2 cases had the same file name, they always matched. However, after 50 files we noticed in some cases where the amount of generated binaries differed (perhaps due to having several contracts within one file being compiled to different binaries). Nevertheless, the total sizes of binary files generated in the two cases were similar, so we thought it was still fair to compare compilation times in this setting [2].

We stopped after collecting 100 files due to the large amount of time needed to check if all these files can be compiled without error, and for reasons like the asynchronous compiling process and the original compiling process each generated a tiny amount of binaries that the other didn't produce.

Ratio async:no-async to compile N different solidity files on x-cpu machine



Results shown above demonstrate expected speedups with 4/8 core google cloud instances, particularly over a large number of files. However, these ratios were smaller than the ones we derived when doing the compiling tests on identical files, which is most likely due to large singular files taking over in one specific processor. These large files probably explain occasional dips in ratio seen, as well as discrepancies for smaller N compared to the previous graph.

We have developed a parallelisation tool inspired by the result of the asynchronous test, see [tutorial/manual](#) for further details.

Nonetheless, as discussed above and in the tutorial, there are several limitations to using this tool, mainly dependency issues.

In order to combat this, we attempted briefly to look at the second parallelisation approach, which is parallelising within the compiler itself. We specifically focused on the following part of code in CompilerStack.cpp:

```
bool CompilerStack::compile()
{
    t_stack.push("CompilerStack::compile");
    if (m_stackState < AnalysisSuccessful)
        if (!parseAndAnalyze()) {
            t_stack.pop();
            return false;
        }

    map<ContractDefinition const*, eth::Assembly const*> compiledContracts;
    for (Source const* source: m_sourceOrder)
        for (ASTPointer<ASTNode> const& node: source->ast->nodes())
            if (auto contract = dynamic_cast<ContractDefinition const*>(node.get()))
                if (isRequestedContract(*contract))
                    compileContract(*contract, compiledContracts);

    this->link();
    m_stackState = CompilationSuccessful;
    t_stack.pop();
    return true;
}
```

Trying to apply async to CompileContract() caused several issues as one cannot apply async on member functions. Using bind was attempted, but unfortunately did not produce good results.

We concluded that parallelisation at lower level would require much more involved and complicated processes; perhaps not achievable (or wanted) within the time frame we had [1].

5. Parallelisation Tool

(See Manual for in detail about every element - located on dev_parallel_tool)

The asynchronous test demonstrated the potential to speed up the compiling process by parallelising it. Therefore, we decided to build a parallelisation tool.

Since some solidity files use the “import” statements, they will have some assumptions about the current directory of the OS when they are being compiled. Therefore, in the original design, we would like to use the C++ filesystem library to set the current path for different tasks in different threads. However, we soon discovered that the parallelisation tool would give a non deterministic output, given the above design. A compiling task may fail in one run, but it may succeed in producing binaries in another run. This is due to that all threads from a process share the same current_path. Thus, it is not possible to have a consistent current_path when compiling in parallel. To tackle this issue we decided to keep the current path as the top level directory, and required users to use absolute path/ relative path instead of imposing any assumption on the current directory. For example, “import x” (this would assume x is in the top

level directory) should be changed to “import ./x” (this would assume x is in the same directory as the file being compiled, which is desired). It solved the problem.

Another thing we need to tackle is the dependency issue in some solidity libraries. We have found that some solidity files that involved operations like name base inheritance, need to be compiled together to generate the binary files. Therefore, it will break the parallelisation process if we separate these files and assign them to different tasks. We address this issue by asking users to create a subdirectory under the top level directory for each batch of interdependent solidity files. For files that are independent and can be compiled in any order, the users can put them in the top level directory. Our helper function `distribute_tasks()` will handle the above two cases by checking its boolean argument `isRoot` when traversing the directories [2].

There are three major functions that support the tool.

```
/* A bare-bones function to compile solidity source files */
int compile(string compiler, string flags, string files,
std::experimental::filesystem::path current);

/* A helper function to distribute sources to different tasks correctly */
void distribute_tasks(string& path, string& extension,
vector<tuple<std::experimental::filesystem::path, vector<string>>>& pool, bool isRoot);

/* A helper function to evenly distribute independent sources (only exist in the
top-level directory according to our assumption) among CPUs */
void parallel_compile(string& compiler, string& flags, vector<string>& sources,
vector<future<int>> &vec_res);
```

And we call them in `main()` in the following way:

```
vector<tuple<std::experimental::filesystem::path, vector<string>>> pool;
distribute_tasks(path, extension, pool, true);
for (auto & p : pool) {
    auto[path, sources] = p;
    if (path == root_path)
        parallel_compile(compiler, flags, sources, vec_res);
    else
```

```

        /* ... */
        vec_res.push_back(async(std::launch::async, compile, compiler, flags,
tasks, path));
    }
    for_each(vec_res.begin(), vec_res.end(), [](future<int> &res){ res.get(); });

```

We first call `distribute_tasks()` to populate the pool vector that stores a tuple of <directory, sol files> for different threads' compiling tasks. Then we iterate over the vector, and if the directory in the tuple is the top level directory, we will call `parallel_compile()` to evenly distribute them to different CPUs. Otherwise we will assume all files in the directory need to be compiled together and we will assign a new thread that compiles all of them.

The parallelisation tool takes 4 arguments, which are path to the solidity compiler executable, flags for the compiler, top level directory for sources and extension for solidity files. Once the parallelised task finishes, the binaries generated will be in the "bin" directories in top level directories or each subdirectory [2].

6. Future Work/Version 1.2:

- Utilize topological sort algorithm to develop parallelisation tool further, searching for imports within files to build a sorted dependence tree structure of sorts, giving less constraints than assuming pure independence
- Report more types of statistics:
 - Compiler memory usage
 - Space on the blockchain
- Use tool to investigate the rest of the codebase
- Modify tool to support concurrency
- Extending parallelization to other single-threaded compilers

7. References

- [1] "ethereum/solidity - GitHub." <https://github.com/ethereum/solidity>. Accessed 26 Apr. 2018.
- [2] "raphael-s-norwitz (Raphael Norwitz) · GitHub." <https://github.com/raphael-s-norwitz>. Accessed 26 Apr. 2018.
- [3] "gcc/timevar.c at master · gcc-mirror/gcc · GitHub." <https://github.com/gcc-mirror/gcc/blob/master/gcc/timevar.c>. Accessed 26 Apr. 2018.
- [4] "c++ - Are system() calls evil? - Stack Overflow." <https://stackoverflow.com/questions/8086071/are-system-calls-evil/8086130>. Accessed 26 Apr. 2018.

- [5] "Solidity by Example — Solidity 0.4.21 documentation - Read the Docs."
<http://solidity.readthedocs.io/en/v0.4.21/solidity-by-example.html>. Accessed 26 Apr. 2018.
- [6] "GitHub - allenday/github-solidity-all: All *.sol files from github in a"
<https://github.com/allenday/github-solidity-all>. Accessed 26 Apr. 2018.

All code taken from our github repository: github.com/raphael-s-norwitz/solidity