

Solidity Compiler Tools User Manual

Killian Rutherford (KRR2125)

Raphael Norwitz (RSN2117)

Jiayang Li (JL4305)

1. Introduction
2. FTime-Report Tool
 - a. Elements
 - i. TimeNode
 - ii. TimeNodeStack
 1. void push(std::string name)
 2. std::string pop()
 3. std::string printString(bool tree)
 4. void print()
 5. void print_recursive(const TimeNode& x, const std::string& arrow, std::stringstream& ss, bool tree)
 - iii. TimeNodeWrapper
 1. Constructor
 2. Destructor
 3. void TimeNodeWrapper::pop()
 - b. Usage
 - i. Flag details
 - ii. Within Compiler Source Code
 - c. Test Cases and Results
3. Attempts at Optimising Using Tool
 - a. Optimise()
 - b. Analyze()
4. Parallelisation
 - a. Usage
 - b. Elements
 - i. Join_contracts
 - ii. Time_test
 - c. Other
 - i. Diff.py
 - ii. Copyshell script
 - d. Results
5. Parallelisation Tool
 - a. Elements
 - i. int compile()
 - ii. void distribute_tasks()
 - iii. void parallel_compile()
 - b. Usage
 - i. Command line details
 - ii. [path_to_compiler]
 - iii. [flags]
 - iv. [directory]
 - v. [extension]
 - c. Test Cases
6. References

1. Introduction

This User Manual goes through the f-time profiling and the parallelism tool we created, diving into specific usage, test cases and function implementations. The manual also documents our attempts at using these tools to optimise the runtime of the solidity compiler. For details on the solidity language, see the official documentation at: <https://solidity.readthedocs.io/en/v0.4.23/>. [1] After discussing with the solidity compiler community, who liked the idea of a built-in profiling tool, we designed our tool based off of Ftime Report in GCC [2]. Having struggled to find concrete optimizations with the profiling tool, we moved on to build parallelized compilation tool, which we rigorously document here.

See the [Design Document](#) for a more linear narrative of the problems we encountered and decisions we made.

For a lighter and more practical introduction to our tool, see the [Tutorial](#).

All code is found on our github page <https://github.com/raphael-s-norwitz/solidity> [3] under various branches (dev_tool, dev_tool_wrapper, dev_optimise, dev_optimise_profile, dev_analyze_profile, dev_junk, dev_parallel_tool, dev_parse_profile).

2. FTime-Report Tool

a. Elements

The below elements can be found on either the dev_tool or dev_tool_wrapper branch. Dev_tool consists of our original design, dev_tool_wrapper includes a wrapper which implements RAI more thoughtfully and handles pop() in the destructor as opposed to users having to call it. [3]

i. TimeNode

The TimeNode class is defined in libsolidity/interface/FTime.h and FTime.cpp files. It's private members consist of chrono::high_resolution_clock times to define begin and end; along with set functions which set these values correspondingly (setBegin(), setEnd()) and get functions which obtain these values (getBegin(), getEnd()).

The constructor simply instantiates the vector of TimeNodes which acts as a list of that TimeNode's children.

No specific destructor is specified - default destructor [3].

ii. TimeNodeStack

TimeNodeStack class is defined in libsolidity/interface/FTime.h and FTime.cpp files. Private members include a vector of TimeNodes serving as the stack, as well as a print_stack vector of TimeNodes and a start time to calibrate all timing differences printed out.

A boolean print_flag is also defined (if this is set to true, pop() will print new nodes to stdout) - which is useful if the --ftime-report flag has been defined. There is also an plain output flag --ftime-report-no-tree.

The Constructor simply instantiates start time.

The Destructor checks if the stack is non empty and gives an error if it is not; otherwise falls to the default destructor [3].

1. void push(std::string name)

Creates a TimeNode object with the given name argument. It then sets the begin time of that node and pushes it onto the stack.

2. std::string pop()

If the stack has multiple elements, the last node from the stack is popped, the end time of it is set, and it is added to the children list of the new top node of the stack.

If the stack has one element, it means the end of the profiling because the pops and pushes for the highest level call has occurred. Thus the time report should be printed with a printString() call, or it can be added to the print_stack, which can be dumped to standard out with an explicit print() call. A pop() call on an empty stack throws an std::runtime_error.

The return value is the string name member of the TimeNode popped [3].

3. std::string printString(bool tree)

Prints all nodes to a string using the print_recursive function. Note that if the print_flag is on, anytime the stack size reaches 1, the time report will be printed to stdout, and printString() will return an empty output. Hence if a user wants to call printString on a local TimeNodeStack, which only looks at elements in the print_stack, the print_flag should remain off, which it is by default. TimeNodes which have yet to be popped will not show up in the string returned.

4. void print()

Simply prints the return value of printString to stdout.

5. void print_recursive(const TimeNode& x, const std::string& arrow, std::stringstream& ss, bool tree)

Recurse through all the node's children and print out in format. If bool tree is true, format the output as a tree structure by adding indents to the arrow string.

iii. TimeNodeWrapper

TimeNodeWrapper class is defined in libsolidity/interface/FTime.h and FTime.cpp files; on the dev_tool_wrapper branch.

We introduced this class so that we could utilize RAI to call pop() in the class destructor, minimizing the number of statements needed to profile a block of code.

Private members of string name, popped flag (in case pop() was called before destructor); and TimeNodeStack reference member [3].

1. Constructor

Takes as input TimeNodeStack (whether a globally or locally defined one), and string name; Calls the TimeNodeStack.push(name), i.e. creates a new node and pushes it on the stack.

2. Destructor

Checks if the popped flag is on; does nothing if it is set, otherwise pop the stack and set the flag to true and default destruct. Note that we are operating under the assumption TimeNodeWrapper objects will be destructed in reverse order they are constructed (put on the Memory Stack Frame) within a function block. If this is violated, the names will not match up and an error will be thrown (see below).

3. void TimeNodeWrapper::pop()

Simply checks by name if the node has already been popped (throws error if it has). Otherwise it pops the TimeNodeStack reference and does an extra check to make sure the popped name is equal to the name saved in the constructor (i.e. the right Node is being popped). It throws an std::runtime_exception if these do not match.

b. Usage

Please see <https://solidity.readthedocs.io/en/v0.4.23/> [1] for solidity documentation and how to build the compiler. Once the compiler solc has been built running the ./scripts/build.sh script on a branch which supports (e.g. dev_tool, dev_tool_wrapper), the flag can be fed into the solc command line arguments.

i. Flag details

Once the solidity compiler has been compiled, and the solc executable produced, a solidity file or multiple solidity files can be compiled through solc A.sol B.sol

Please see the solidity documentation for explanation regarding already existing flags, such as --optimize and --bin to produce binary source files. The FTimeReport flag is added with --ftime-report or --ftime-report-no-tree

The no-tree flag prints what has been profiled without the tree indentation, which may be desirable in cases where there are many nested profiling statements.

The flag is handled in solc/CommandLineInterface.cpp in the handleFTimeReport function. This only deals with printing the global t_stack variable which is defined in FTime.h; by printing everything in the print_stack and setting the print_flag on for future additions [3].

ii. Within Compiler Source Code

To profile code within the compiler, the global t_stack variable, which has been defined in FTime.h, should be used. The print() functions and print_flags are called in the CommandLineInterface.cpp file, so no additional work is needed to add debug statements to the existing output.

Local TimeNodeStacks can also be defined and used, but in this case the print_flag of the local TimeNodeStack needs to be turned on when the stack is defined or print() must be called to get the debug string. Note that the TimeNodeStack needs to be accessible everywhere you wish to put profiling statements.

The FTime-report tool source is located in libsolidity/interface, so to import it into other libraries within solidity, it is necessary to add the libsolidity to the target_link_libraries within the specific library's CMakeLists.txt [4].

Besides this, one needs to add #include <libsolidity/interface/FTime.h> at the beginning of the source file in which the code wanting to be profiled is located.

Then, one can either explicitly stack.push(string) and stack.pop() around the code block wanting to be profiled, or use the RAIL TimeNodeWrapper. See tutorial for further details.

c. Test Cases and Results

See tutorial for simple test cases; as well as the design document and sections below for further uses. Boost test cases are located on the dev_tool or dev_tool_wrapper branch; in the test/libsolidity/FTimeTest.cpp file.

Some of the tests are simple, and check that standard functions like push() and pop() don't throw errors when and only when expected. We use std::regex to evaluate the printed output, making sure that the trees contained the correct strings and were indented as expected. The most complicated tests involved string parsing with regex; as these were quite challenging to get right. Please see the above file again for specific cases.

3. Attempts at Optimising Using Tool

Initial runs of the FTimeReport Tool show optimise, parse and analyse functions within the compiler taking a lot of time.

a. Optimise()

Using the tool on a simple contract Coin.sol, we see CompilerContext::optimise takes up quite a lot of the time. Below results are conducted on the dev_optimise_profile branch.

```
killianrutherford1@mbp ~/Desktop/Courses/cppDesign/proj/solidity $ solc --ftime-report Coin.sol
```

Warning: This is a pre-release compiler version, please do not use it in production.

Coin.sol:15:5: Warning: Defining constructors as functions with the same name as the contract is deprecated. Use "constructor(...) { ... }" instead.

```
function Coin() public {
```

^ (Relevant source part starts here and spans across multiple lines).

namespace/function name	unix begin time(µs)	time elapsed(µs)
CLL::readInputFilesAndConfigureRemappings	285	221
CompilerStack::setRemappings	516	0
CompilerStack::setEVMVersion	554	0
CompilerStack::compile	557	3629
_CompilerStack::parse	557	125
_CompilerStack::analyze	683	1247
_CompilerStack::compileContract: Coin	1934	2239
_CompilerStack::createMetadata	1942	223
_Compiler::compileContract	2272	1582
_ContractCompiler::compileContract	2277	400
_ContractCompiler::initializeContext	2283	20
_ContractCompiler::appendFunctionSelector	2304	234
_ContractCompiler::appendMissingFunctions	2542	134
_ContractCompiler::initializeContext	2684	7
_ContractCompiler::appendMissingFunctions	2722	0

_ContractCompiler::appendMissingFunctions	2724	0
_CompilerContext::optimise	2726	1115
_Assembly::optimise	2728	1103
_OptimiserSettings::optimise	2728	1093
_OptimiseInternal	2729	1082
_OptimiseInternal	2732	939
_test1	2734	10
_runPeephole	2744	615
_PeepholeOptimiser::optimise	2746	297
_while loop ApplyMethods	2747	288
_PeepholeOptimiser::optimise	3045	307
_while loop ApplyMethods	3045	279
_runDeduplicate	3361	0
_test3	3362	0
_test1	3362	8
_runPeephole	3371	298
_PeepholeOptimiser::optimise	3371	291
_while loop ApplyMethods	3371	265
_runDeduplicate	3670	0
_test3	3670	0
_test1	3675	1
_runPeephole	3677	58
_PeepholeOptimiser::optimise	3677	26
_while loop ApplyMethods	3677	24
_PeepholeOptimiser::optimise	3704	28
_while loop ApplyMethods	3704	25
_runDeduplicate	3737	0
_test3	3737	0
_test1	3737	0
_runPeephole	3738	32
_PeepholeOptimiser::optimise	3738	30
_while loop ApplyMethods	3739	27
_runDeduplicate	3771	0
_test3	3771	0
_test1	3772	0
_runPeephole	3772	33
_PeepholeOptimiser::optimise	3773	31
_while loop ApplyMethods	3773	28
_runDeduplicate	3806	0
_test3	3806	0
_ContractCompiler::initializeContext	3933	6
_ContractCompiler::appendMissingFunctions	3982	0
_Assembly::optimise	3983	156
_OptimiserSettings::optimise	3983	147
_OptimiseInternal	3983	138
_OptimiseInternal	3984	8
_test1	3984	0
_runPeephole	3984	3
_PeepholeOptimiser::optimise	3984	1
_while loop ApplyMethods	3985	0

_runDeduplicate	3988	0
_test3	3989	0
_OptimiseInternal	3993	23
_test1	3993	0
_runPeephole	3994	20
_PeepholeOptimiser::optimise	3994	18
_while loop ApplyMethods	3994	15
_runDeduplicate	4015	0
_test3	4015	0
_test1	4017	1
_runPeephole	4019	63
_PeepholeOptimiser::optimise	4019	28
_while loop ApplyMethods	4020	24
_PeepholeOptimiser::optimise	4048	31
_while loop ApplyMethods	4049	28
_runDeduplicate	4084	0
_test3	4084	0
_test1	4084	0
_runPeephole	4085	31
_PeepholeOptimiser::optimise	4085	29
_while loop ApplyMethods	4086	26
_runDeduplicate	4117	0
_test3	4118	0
CommandLineInterface::actOnInput	4243	1018

Using the tool in further detail, we break up into further levels of detail above, and see that almost all of the time in `CompilerContext::optimise` is split up in the while loop `applyMethods` in `libevmasm/PeepholeOptimiser.cpp` (see snippet below). [4]

```
bool PeepholeOptimiser::optimise()
{
    t_stack.push("PeepholeOptimiser::optimise");
    OptimiserState state {m_items, 0, std::back_inserter(m_optimisedItems)};
    t_stack.push("while loop applyMethods");
    while (state.i < m_items.size())
        applyMethods(state, PushPop(), OpPop(), DoublePush(), DoubleSwap(), CommutativeSwap(), SwapComparison(), JumpToNext(), UnreachableCode(), TagConjunctions(), Identity());
    t_stack.pop();
    if (m_optimisedItems.size() < m_items.size() || (
        m_optimisedItems.size() == m_items.size() && (
            eth::bytesRequired(m_optimisedItems, 3) < eth::bytesRequired(m_items, 3) ||
            numberOfPops(m_optimisedItems) > numberOfPops(m_items)
        ))
    )
    {
        m_items = std::move(m_optimisedItems);
        t_stack.pop();
        return true;
    }
    else {
        t_stack.pop();
        return false;
    }
}
```

As seen below, there's no one `applyMethods` that seems to be taking that much longer than any other (i.e. the methods). It seems rather a combination of all of them which takes a long time, so if their calls can be minimized somehow, that may improve performance. See Design document for some attempts that were made.

_CompilerContext::optimise	2739	50342
_Assembly::optimise	2741	46939
_OptimiserSettings::optimise	2741	43510
_OptimiseInternal	2741	40108
_OptimiseInternal	2743	33196
_test1	2745	9
_runPeephole	2755	18447
_PeepholeOptimiser::optimise	2756	7466
_while loop ApplyMethods	2757	6473
_applyMethods	2757	17
_applyMethods	2757	14
_applyMethods	2757	12
_applyMethods	2757	10
_applyMethods	2758	8
_applyMethods	2758	6
_applyMethods	2759	4
_applyMethods	2760	2
_applyMethods	2761	0
_applyMethods	2761	0
_applyMethods	2775	16
_applyMethods	2776	13
_applyMethods	2776	9
_applyMethods	2776	7
_applyMethods	2776	5
_applyMethods	2776	3
_applyMethods	2777	2
_applyMethods	2777	1
_applyMethods	2777	0
_applyMethods	2777	0
_applyMethods	2793	13
_applyMethods	2793	11
_applyMethods	2793	9
_applyMethods	2793	7
_applyMethods	2793	5
_applyMethods	2793	3
_applyMethods	2794	2
_applyMethods	2794	1
_applyMethods	2794	0
_applyMethods	2795	0
_applyMethods	2809	13
_applyMethods	2809	11
_applyMethods	2809	8
_applyMethods	2809	6
_applyMethods	2809	5
_applyMethods	2809	3
_applyMethods	2810	1
_applyMethods	2810	1
_applyMethods	2810	0
_applyMethods	2810	0
_applyMethods	2823	13

_applyMethods	2824	11
_applyMethods	2824	8
_applyMethods	2824	6
_applyMethods	2824	5
_applyMethods	2824	3
_applyMethods	2824	2
_applyMethods	2825	1
_applyMethods	2825	0
_applyMethods	2825	0

Even with larger contracts (compilationTests/corion/provider.sol), we see optimise and applymethods taking large amounts of time, so this seems like a good place to be looking to shave off compile time. [4]

```
killianrutherford1@mbp ~/Desktop/Courses/cppDesign/proj/solidity $ solc --ftime-report
test/compilationTests/corion/provider.sol | grep ApplyMethods
```

_while loop ApplyMethods	242722	11
_while loop ApplyMethods	248394	3557
_while loop ApplyMethods	252041	3350
_while loop ApplyMethods	255873	3300
_while loop ApplyMethods	259627	3523
_while loop ApplyMethods	263523	467
_while loop ApplyMethods	264004	455
_while loop ApplyMethods	264471	452
_while loop ApplyMethods	264989	452
_while loop ApplyMethods	266536	0
_while loop ApplyMethods	266547	14
_while loop ApplyMethods	266587	471
_while loop ApplyMethods	267071	458
_while loop ApplyMethods	267540	457
_while loop ApplyMethods	268063	459
_while loop ApplyMethods	271548	290
_while loop ApplyMethods	271847	289
_while loop ApplyMethods	272176	288
_while loop ApplyMethods	272496	11
_while loop ApplyMethods	272509	11
_while loop ApplyMethods	272527	11
_while loop ApplyMethods	272543	11
_while loop ApplyMethods	272675	0
_while loop ApplyMethods	272699	13
_while loop ApplyMethods	272729	10
_while loop ApplyMethods	272742	11
_while loop ApplyMethods	272758	10
_while loop ApplyMethods	283165	7557
_while loop ApplyMethods	290915	6169
_while loop ApplyMethods	297218	5768

_while loop ApplyMethods	303377	5787
_while loop ApplyMethods	309959	7950
_while loop ApplyMethods	318073	6675
_while loop ApplyMethods	325557	6284
_while loop ApplyMethods	332443	142

b. Analyse()

According to the ftime report, the analyze pass CompilerStack::analyze generally takes around 50% of the total compile time when compiling the Coin contract. Below results are conducted on the dev_analyze_profile branch.

namespace/function name	unix begin time(μs)	time elapsed(μs)
CommandLineInterface::processInput()	358	3384
_CLI::readInputFilesAndConfigureRemappings	363	257
_CompilerStack::setRemappings	631	2
_CompilerStack::setEVMVersion	680	0
_CompilerStack::compile	684	3016
_CompilerStack::parse	684	175
_CompilerStack::analyze	860	1573
_CompilerStack::compileContract: Coin	2438	1254
_CompilerStack::createMetadata	2450	266
_Compiler::compileContract	2856	661
_ContractCompiler::compileContract	2861	303
_ContractCompiler::initializeContext	2869	26
_ContractCompiler::appendFunctionSelector	2896	94
_ContractCompiler::appendMissingFunctions	2994	168
_ContractCompiler::initializeContext	3175	7
_ContractCompiler::appendMissingFunctions	3200	0
_ContractCompiler::appendMissingFunctions	3201	0
_CompilerContext::optmise	3203	311
_ContractCompiler::initializeContext	3603	6
_ContractCompiler::appendMissingFunctions	3641	0
CommandLineInterface::actOnInput	3757	172

Splitting the analyze pass further by applying our tools to .sol files that is relevant to the analyze pass, we have the following result:

namespace/function name	unix begin time(μs)	time elapsed(μs)
CLI::readInputFilesAndConfigureRemappings	376	263
CompilerStack::setRemappings	650	2
CompilerStack::setEVMVersion	698	0
CompilerStack::compile	701	4253
_CompilerStack::parse	701	167
_CompilerStack::analyze	872	1588
_SyntaxChecker::checkSyntax	874	31
_SemVerMatchExpressionParser::parse	882	4
_SemVerMatchExpression::matches	888	0
_DocStringAnalyser::analyseDocStrings	907	16
_GlobalContext::declarations	1036	1

_NameAndTypeResolver::NameAndTypeResolver	1040	20
_GlobalContext::setCurrentContract	1095	0
_GlobalContext::currentThis	1096	1
_GlobalContext::currentSuper	1101	1
_ReferencesResolver::resolve	1110	7
_ReferencesResolver::resolve	1117	5
_ReferencesResolver::resolve	1125	1
_ReferencesResolver::resolve	1127	0
_ReferencesResolver::resolve	1129	4
_ReferencesResolver::resolve	1134	1
_ReferencesResolver::resolve	1136	0
_ReferencesResolver::resolve	1137	1
_ReferencesResolver::resolve	1139	0
_ReferencesResolver::resolve	1140	5
_ReferencesResolver::resolve	1151	4
_ReferencesResolver::resolve	1156	6
_TypeChecker::checkTypeRequirements	1164	216
_PostTypeChecker::check	1382	4
_StaticAnalyzer::analyze	1388	15
_ViewPureChecker::check	1404	24
_SMTChecker::SMTChecker	1429	831
_SMTChecker::analyze	2261	15
_CompilerStack::compileContract: Coin	2469	2478
_CompilerStack::createMetadata	2482	323
_Compiler::compileContract	2928	1739
_ContractCompiler::compileContract	2934	510
_ContractCompiler::initializeContext	2941	26
_ContractCompiler::appendFunctionSelector	2968	287
_ContractCompiler::appendMissingFunctions	3260	182
_ContractCompiler::initializeContext	3451	7
_ContractCompiler::appendMissingFunctions	3499	0
_ContractCompiler::appendMissingFunctions	3500	0
_CompilerContext::optimise	3503	1162
_ContractCompiler::initializeContext	4752	8
_ContractCompiler::appendMissingFunctions	4813	0
CommandLineInterface::actOnInput	5069	234

As you can see, SMTChecker::SMTChecker and TypeChecker::checkTypeRequirements took most of the time in the analyze pass. For SMTChecker, the compiler spent more than 800 microseconds on the constructor of the object [3]. The relevant part of the code follows:

```

if (noErrors)
{
    t_stack.push("SMTChecker::SMTChecker");
    SMTChecker smtChecker(m_errorReporter, m_smtQuery);
    t_stack.pop();
    for (Source const* source: m_sourceOrder)
        smtChecker.analyze(*source->ast);
}

```

libsolidity/interface/CompilerStack.cpp

Next, let's look at the SMTChecker constructor in libsolidity/formal/SMTChecker.cpp

```
SMTChecker::SMTChecker(ErrorReporter& _errorReporter, ReadCallback::Callback const& _readFileCallback):
#ifdef HAVE_Z3
    m_interface(make_shared<smt::Z3Interface>()),
#elif HAVE_CVC4
    m_interface(make_shared<smt::CVC4Interface>()),
#else
    m_interface(make_shared<smt::SMTLib2Interface>(_readFileCallback)),
#endif
    m_errorReporter(_errorReporter)
{
    (void)_readFileCallback;
}
```

Another thing we tried was to apply our profiling tool to Z3Interface constructor, and it turned out that when the Z3Interface constructor was called, approximately 800 microseconds already passed since ViewPureChecker::check, which is the previous pass of SMTChecker constructor, finished. Therefore, we believe the std::function ReadCallback::Callback const& _readFileCallback is a major bottleneck of the compiler's speed.

In terms of TypeChecker::checkTypeRequirements, we decided to divide the checkTypeRequirements function into 3 parts and recorded the time spent on each part. We found that the third part is what usually makes the type checking slow. [4]

```
bool TypeChecker::visit(ContractDefinition const& _contract)
{
    t_stack.push("TypeChecker::visit 1");
    m_scope = &_contract;
    /* omitted */
    t_stack.pop();
    t_stack.push("TypeChecker::visit 2");
    /* omitted */
    t_stack.pop();
    t_stack.push("TypeChecker::visit 3");
    for (auto const& n: _contract.subNodes())
        if (!visited.count(n.get()))
            n->accept(*this);
    t_stack.pop();
    /* omitted */
    return false;
}
```

libsolidity/analysis/TypeChekcer.cpp

_TypeChecker::checkTypeRequirements	1233	228	
_TypeChecker::visit 1	1234	49	
_TypeChecker::checkContractDuplicateFunctions	1237	11	
_TypeChecker::checkContractDuplicateEvents		1250	3
_TypeChecker::checkContractIllegalOverrides	1254	4	
_TypeChecker::checkContractAbstractFunctions	1260	10	
_TypeChecker::checkContractBaseConstructorArguments	1276	4	
_TypeChecker::visit 2	1284	1	
_TypeChecker::visit 3	1287	100	
_TypeChecker::checkContractExternalTypeClashes	1388	38	

Ftime-report targeting TypeChecker

The for loop in TypeChecker::visit 3 takes around 50% of the time of the TypeChecker::checkTypeRequirements pass. In this piece of the code, the compiler is visiting each ASTNode in the contract and checking if there is any error. Therefore, we believe it is reasonable for this part to use longer time and it would be difficult, if not impossible, to speed this part up.

4. Parallelisation

For details into the thought processes behind parallelisation and methods attempted, please see the Design Document section 3.

The final parallelisation code which was used to run results is located in the dev_optimise branch, in the directory parallelisation_test/ [3].

a. Usage

This code is to see the effect of using C++'s async and seeing possible time gains to running several number of source files within a library split over different numbers of processors.

E.g. given 10 solidity files, A.sol - J.sol, and 4 processors, we would compare the times running:

```
solc A.sol B.sol C.sol D.sol E.sol F.sol G.sol H.sol I.sol J.sol
(on 1 processor)
Vs.
```

```
CPU1: solc A.sol B.sol C.sol
CPU2: solc D.sol E.sol F.sol
CPU3: solc G.sol H.sol
```

CPU4: solc I.sol J.sol

As is mentioned, there are several limitations to this which the user should be aware of by reading the design document, mainly that these files must be independent from one another, and cannot have interlinked dependencies or redefinitions of classes, etc...

The Makefile specifies C++ 17 to be used, as experimental filesystem is needed to traverse filesystems and directories. (In the future this may need to be changed to just filesystem)

Compiling the executable can be done by simply typing make within the parallelisation_test/ directory. A test executable is produced, and is run with two command line arguments: N (the number of files wanting to run the comparison over), and the path to the directory containing these files. The source code described below is mainly located in parallelisation_test/select_contracts_2.cpp [3].

b. Elements

i. Join_contracts

In order to find out if N number of files from a directory path given can be compiled with one another; and potentially parallelised over multiple CPU's, join_contracts takes in a compiler string, a select_flag for this compiler, a path to the directory, a test string consisting of the files to compile and N.

It iterates through the solidity files in the directory using experimental::filesystem. The function adds the new file to the test_string, and runs the compiler with the test_string, checking to see if the binary produced is valid.

If it is, the test_string is updated to include this new file. The above process is repeated and continues until N is reached or until all files in the directory are exhausted.

Join_contracts returns the number of files it was able to compile, if this is less than the original N, this means not enough files in the directory were able to satisfy the user's initial request.

ii. Time_test

After join_contracts is executed, and the test_string variable contains N files which do not have interdependencies or issued, time_test is called; taking in arguments of vector of all the files, test_string, compiler command and flags, and directories for single thread and multi thread binary output.

Using the thread library, the number of CPU's on the machine is calculated, the files are split optimally into different strings for each CPU. For multiprocessing, a results vector of `future<int>` is created. Using the `async()` call, each task on a CPU is executed and pushed onto the vector. `get()` is later on called on each element of the vector to obtain the timing result. This is done using `chrono::high_resolution` clock; and single processor compiling is also executed, printing out both results.

All binary files during compilation are outputted to respective single and multi processing directories.

See design document for some notes on the binary files produced in some cases with large numbers of varied files.

c. Other

i. Diff.py

A python script `diff.py` is also provided in the `parallelisation_test/` directory to ensure all binary files within both single thread and multi thread directories match.

ii. Copyshell script

A bash shell script `copy_shell.sh` is provided to automatically run three consecutive times: make the test executable, run with a specific N and path string, and run the diff python file on the the binary output directories.

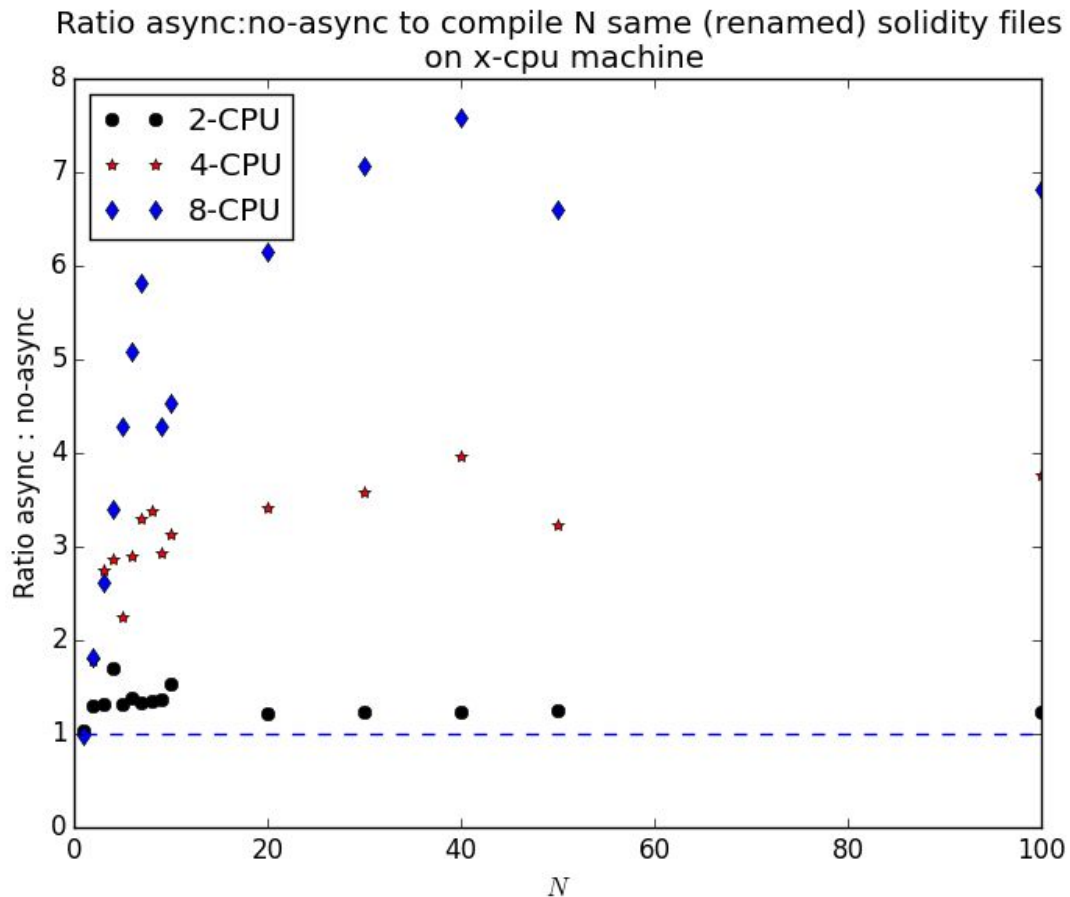
d. Results

(See Design Document for more in-depth discussion)

The first dataset we tested on was selecting a random contract, and copying it a certain number of times, whilst altering the name of the contract classes and functions, to prevent as much as possible the compiler from optimizing and predicting / cache hits, thereby trying to force it to recompile every single contract.

This was done with `BlindAuction.sol` file, which contains both a `BlindAuction` contract and a purchase contract. (see <https://solidity.readthedocs.io/en/v0.4.23/solidity-by-example.html> for the files) [1].


Time ratios were averaged over 3 runs and results over different number of google instance n1-standard vCPU are shown. (n1-standard-2 (2 vCPUs, 7.5 GB memory) ; n1-standard-4 (4 vCPUs, 15 GB memory) ; n1-standard-8 (8 vCPUs, 30 GB memory) [5].



We also wanted to test out parallelisation on a more realistic data set (i.e. not just copies of a specific file) to see what could be a potential “real-world” speedup to benchmark and to avoid any thread taking advantage of cache memory. We were pointed by the solidity developer’s team towards <https://github.com/allenday/github-solidity-all>, [6] which contains all solidity files on github collected in one place.

We tested on 100 files from this repo (See Design Document):
8894 lines; 33,362 words; 294,515 characters

Ratio async:no-async to compile N different solidity files on x-cpu machine



5. Parallelisation Tool

The tool should be accessed in the `dev_parallel_tool` branch in the `parallel/` directory.

a. Elements

i. `int compile(string compiler, string flags, string files, std::experimental::filesystem::path current)`

The `compile()` function is defined in `parallel/cmp.cpp`. It is a function that generates the command to compile the solidity files and then calls the host environment's command processor to execute the command.

ii. void distribute_tasks(string& path, string& extension, vector<tuple<std::experimental::filesystem::path, vector<string>>>& pool, bool isRoot)

The `distribute_tasks()` function is defined in `parallel/cmp.cpp`. It populates the `pool` vector by filling it with batches of files that will be compiled together. It will leverage the C++ experimental filesystem library to search for files to compile. When `isRoot` is true, the function will consider all solidity files in the **path** to be independent from each other, and thus they can be distributed to different threads in any order. When `isRoot` is false, it will distribute all solidity files in **path** to a single task since we assume any solidity files in subdirectories need to be compiled together [3].

ii. void parallel_compile(string& compiler, string& flags, vector<string> &sources, vector<future<int>> &vec_res)

The `parallel_compile()` function is defined in `parallel/cmp.cpp`. It will evenly partition files in **sources** into different batches for compiling, and it will consider the hardware resources of the machine for the partition. Each partition will then be passed to a **compile()** task and then added to the **vec_res** waiting to be executed parallelly [3].

b. Usage

i. Command line details

To compile the parallelisation tool, use gcc 7 and run the following command:

```
G++-7 -o parallel -std=c++17 -O2 par_cmp.cpp -lstdc++fs -lpthread
```

The **parallel** executable takes 4 arguments, which are **[path_to_compiler]**, **[flags]**, **[directory]**, and **[extension]**.

An example command:

```
./parallel "solc" "--bin --ignore-missing" "parallel/test" ".sol"
```

ii. [path_to_compiler]

Path to the compiler executable. If the **solc** executable is in `/usr/bin`, then set it to be `"/usr/bin/solc"`. If your environment has **solc** as a global executable, then you can just make the `path_to_compiler` `"solc"`.

iii. [flags]

Flags the users would like to specify for the solidity compiler. If a user just wants the compiler to export binary files, then just use "--bin --ignore-missing". Additionally, since our tool will automatically put the generated binary files under a "bin" folder in each directory, users should not include any output directory in the flags, which means they should not include -o in the flags.

iv. [directory]

The top level directory for the files to be compiled. Put all files that can be compiled independently under this directory and all other files into subdirectories under this directory.

v. [extension]

The extension of the solidity files. Usually it will just be ".sol". Remember to put a dot before the extension. The parallelisation tool will only compile a file if it finds the extension in its filename.

c. Test Cases

Checkout the dev_parallel_tool branch of our github repo. In the folder **parallel**, there is a Makefile and a **test** directory that contains a collection of solidity files. Generate the parallelisation tool by "make" command and use **test**'s path as the **[directory]** argument for the parallelisation tool and compile the files parallelly. After the compiling process finished, you will find the generated binaries in the bin folders under **test** or any of its subdirectories.

6. References

- [1] "Solidity — Solidity 0.4.23 documentation - Read the Docs." <https://solidity.readthedocs.io/>. Accessed 26 Apr. 2018.
- [2] "Using the GNU Compiler Collection (GCC): Developer Options." <https://gcc.gnu.org/onlinedocs/gcc/Developer-Options.html>. Accessed 26 Apr. 2018.
- [3] "raphael-s-norwitz (Raphael Norwitz) · GitHub." <https://github.com/raphael-s-norwitz>. Accessed 26 Apr. 2018.
- [4] "ethereum/solidity - GitHub." <https://github.com/ethereum/solidity>. Accessed 26 Apr. 2018.
- [5] "Virtual Machine Instances | Compute" <https://cloud.google.com/compute/docs/instances/>. Accessed 26 Apr. 2018.
- [6] "GitHub - allenday/github-solidity-all: All *.sol files from github in a" <https://github.com/allenday/github-solidity-all>. Accessed 26 Apr. 2018.

All code taken from our github repository: github.com/raphael-s-norwitz/solidity [3]
Manual, tutorial and design document located on submissions branch.